

CS 112

Introduction to Data Structures

WEEK 05

Generic Containers

Eric Araújo

Calvin University · Fall 2026

This Week

- 1 The problem with `typedef`
- 2 Class templates
- 3 Turning a class into a template
- 4 STL `vector`
- 5 Iterators
- 6 How long do operations take?
- 7 Amortized constant time **KEY**

One Class, One Type

```
typedef double Item;           // in PyList.h

PyList scores;                 // holds doubles ✓
PyList names;                  // also doubles x
```

Change the `typedef` and every `PyList` in the program changes with it. Want a `PyList` of doubles and a `PyList` of strings in one program? You'd need two copies of the class with different names.

Python never had this problem, a list holds anything. C++ needs types at compile time, so it needs another way.

A Class That Takes a Type

```
template <typename Item>
class PyList {
public:
    void append(const Item &it);
private:
    Item      *myArray;
    unsigned  mySize;
};
```

```
PyList<double>  scores;    // a PyList of doubles
PyList<string>  names;     // a PyList of strings
```

You instantiate a *class* to get an object. You instantiate a *class template* to get a class. `PyList<double>` and `PyList<string>` are two different, separately compiled classes.

Turning a Class Into a Template

- Build and debug it first with `typedef double Item;`: get it working
- Replace the typedef with `template <typename Item>` above the class
- In the `.cpp`, put that same line above **every** method
- Change each `PyList::` to `PyList<Item>::`
- Move the implementation into the header, then `#include` it

```
template <typename Item>
void PyList<Item>::append(const Item &it) { ... }
```

Debug once as a concrete class, then generalise. Templated compiler errors are famously unfriendly, don't meet them and your own bugs at the same time.

TALK TO YOUR NEIGHBOR · TTYN

Why does a class template's implementation go in the header?

- A.** Style, the committee preferred it that way
- B.** The compiler must see the code to generate a class for each type used
- C.** Headers compile faster than source files
- D.** So the linker can find it

You Already Have One: **vector**

```
#include <vector>

vector<int> scores;

scores.push_back(90);           // append
scores.push_back(85);

cout << scores.size();         // 2
cout << scores[0];             // 90
```

The STL is a library of class templates: **vector**, **list**, **stack**, **queue**, **set**, **map**. You've just built a simplified **vector**: from here on you may use the real one.

Iterators

```
vector<string>::iterator it;

for (it = names.begin(); it != names.end(); ++it) {
    cout << *it << endl;    // dereference, like a pointer
}

for (const string &name : names) {    // the modern way
    cout << name << endl;
}
```

- `begin()`: points at the first element
- `end()`: points *just past* the last one
- Every STL container is walked the same way

An iterator behaves like a pointer on purpose: `*it` to read, `++it` to advance.

Counting the Work: Indexing

```
v[i] → *(myArray + i * sizeof(Item))
```

- one multiply, one add, one dereference
- three operations: **whatever i is**
- whatever the size of the array is

Work that doesn't grow with n is **constant time**.
Indexing an array is the classic example.

Counting the Work: Append

Room left

```
myArray[mySize] = it;  
++mySize;
```

2 operations → **constant**

Full

```
allocate 2n  
copy n items  
delete old
```

$\approx 3n + 5 \rightarrow$ **linear**

So which is it? The honest answer is "usually cheap, occasionally expensive", and we need a way to talk about that.

Amortized Constant Time

Start empty and append 16 items, doubling each time we fill up:

Append #	Capacity before	Copies
1	0	0
2	1	1
3	2	2
5	4	4
9	8	8
the other 11	none	0

15 copies across 16 appends, under one copy each *on average*. Spread over the whole sequence, appending is **amortized constant time**: rarely expensive, cheap in the long run.

TALK TO YOUR NEIGHBOR · TTYN

What is the time complexity of `append()` on a dynamic array?

- A. Constant; it's just two operations
- B. Linear; it might copy everything
- C. Amortized constant, usually two operations, occasionally n
- D. It depends on the Item type

TALK TO YOUR NEIGHBOR · TTYN

Inserting at the *front* of a dynamic array is...

- A.** Constant time
- B.** Amortized constant time
- C.** Linear time, every element must shift up one
- D.** Faster than appending

Demo



One template, many types.

(the template compiles fine. Item does not.)

Cost So Far

Operation	Dynamic array	Why
<code>v[i]</code>	constant	address arithmetic
<code>append()</code>	amortized constant	doubling spreads the copy
<code>insert()</code> at front	linear	everything shifts
<code>remove()</code> from middle	linear	everything after shifts
copy constructor	linear	copies every element
<code>getSize()</code>	constant	we stored the count

Keep this table. Next week we build a structure with almost the opposite profile, and the comparison is the point.

Week 05 Recap

- A **class template** takes a type the way a function takes an argument
- `PyList<int>` and `PyList<string>` are two separately generated classes
- Template implementations live in the **header**
- Every STL container is walked with `begin()` and `end()`
- Indexing is constant; append is **amortized constant**; inserting at the front is linear

Next: a06 →

This deck stands on earlier CS112 materials by
Joel Adams and **Victor Norman** · adapted and extended by **Eric Araújo**

