

CS 112

Introduction to Data Structures

WEEK 14

Hash Tables and Sorting

Eric Araújo

Calvin University · Fall 2026

This Week

- 1 Computing where an item lives
- 2 Hash functions
- 3 Collisions **KEY**
- 4 Load factor and the time-space trade
- 5 Selection sort
- 6 Insertion sort
- 7 Where the course has taken you

What If We Just Computed the Index?

A tree finds an item in $O(\lg n)$ by comparing. An array finds index `i` in $O(1)$ by arithmetic.

So: turn the *key itself* into an index. No comparisons, no walking — one calculation and you're there. That's a **hash table**.

A Hash Function

```
unsigned hash(const string &key, unsigned capacity) {  
    unsigned sum = 0;  
    for (char c : key) {  
        sum = sum * 31 + c;           // mix the characters  
    }  
    return sum % capacity;           // fold into the table  
}
```

- **Deterministic** — same key, same slot, every time
- **Fast** — otherwise you've lost the advantage
- **Spreads out** — keys should scatter across the table

The `% capacity` is what forces an unbounded key space into a fixed number of slots — and that is exactly why collisions are unavoidable.

Collisions

slot 0		—
slot 1	"ada" → "bob"	both hashed to 1
slot 2		"grace"

Two keys, one slot. With 4 billion possible strings and 100 slots, this is not bad luck — it's arithmetic.

Chaining — each slot holds a linked list of everything that landed there

Open addressing — on a clash, probe forward for the next free slot

With chaining, what is the worst case for find?

- A.** $O(1)$ — hashing is always constant
- B.** $O(\lg n)$ — the chain is a tree
- C.** $O(n)$ — every key could hash to the same slot
- D.** $O(n^2)$

Load Factor

```
load factor = items / slots
```

- Low (say 0.5) — short chains, fast lookups, memory sitting empty
- High (say 5.0) — memory well used, chains long, lookups slower
- Most implementations **rehash** — double the table and redistribute — around 0.75

This is the **time-space trade-off** in its purest form: buy speed with empty slots. Rehashing is $O(n)$, amortised across the inserts — exactly the doubling argument from week 5.

Where Hash Tables Sit

Operation	Hash table	Balanced tree	Sorted array
find	O(1) avg	$O(\lg n)$	$O(\lg n)$
insert	O(1) avg	$O(\lg n)$	$O(n)$
remove	O(1) avg	$O(\lg n)$	$O(n)$
sorted listing	$O(n \lg n)$	O(n)	O(n)
worst case	$O(n)$	$O(\lg n)$	$O(\lg n)$

Fastest on average, and the only one with no sense of order. Every row of this table is a design decision you can now make on purpose.

Selection Sort

```
for (int i = 0; i < n - 1; ++i) {  
    int smallest = i;  
    for (int j = i + 1; j < n; ++j) {  
        if (a[j] < a[smallest]) { smallest = j; }  
    }  
    swap(a[i], a[smallest]);  
}
```

Find the smallest of what's left, swap it into place, repeat.

Nested loops over $n \rightarrow \mathbf{O(n^2)}$, always. It does the same work on sorted input as on random input, but it makes at most $n-1$ swaps.

Insertion Sort

```
for (int i = 1; i < n; ++i) {  
    Item key = a[i];  
    int j = i - 1;  
    while (j >= 0 && a[j] > key) {  
        a[j + 1] = a[j];    // shift right  
        --j;  
    }  
    a[j + 1] = key;        // drop it in  
}
```

Take the next item, slide it back until it's in place — the way you sort a hand of cards.

Worst case $O(n^2)$, but on **already-sorted** input the inner loop never runs: $O(n)$.

Your data is almost sorted already. Which of the two is better?


- A.** Selection sort — it always does the same work
- B.** Insertion sort — it approaches $O(n)$ when little is out of place
- C.** Identical — both are $O(n^2)$
- D.** Neither works on partly sorted data

Comparing the Two

	Selection sort	Insertion sort
Best case	$O(n^2)$	$O(n)$
Worst case	$O(n^2)$	$O(n^2)$
Swaps / writes	$O(n)$	$O(n^2)$
Stable?	no	yes

Neither is what you'd ship — `std::sort` is $O(n \lg n)$. But every fast sort is built from ideas you can now read, and both of these are the right choice for a small enough n .

Demo



Racing selection sort against std::sort.

(n = 50,000, and then we wait)

Fourteen Weeks

Structure	Good at	Bad at
Array / vector	indexing, appending	inserting at the front
Linked list	inserting and removing anywhere you already are	indexing
Stack / queue	one disciplined end (or two)	looking in the middle
BST	ordered search — when balanced	sorted input, unless balanced
AVL tree	guaranteed $O(\lg n)$	more code, more memory
Hash table	raw lookup speed	any question about order

You built every one of these. The skill you're leaving with isn't any single structure — it's asking *which operations will this code do a million times?* before choosing.

Week 14 Recap

- A hash function turns a key into an index — $O(1)$ on average
- Collisions are inevitable; **chaining** and **open addressing** handle them
- Load factor trades memory for speed; rehashing is the doubling trick again
- Selection sort is $O(n^2)$ always; insertion sort is $O(n)$ on nearly-sorted data
- Choose the structure by the operations you repeat most

Next: Final Window · book by Dec 9 →

This deck stands on earlier CS112 materials by
Joel Adams and **Victor Norman** · adapted and extended by **Eric Araújo**

